

# Database Systems

---

COMP 3010E FALL 2025

LECTURE 14 CONCURRENCY CONTROL

# Agenda

---

- Transaction
- Concurrency Control
  - Lock-Based Concurrency Control Protocols
  - Timestamp-Based Concurrency Control Protocols

# Transaction

---

A transaction is a sequence of actions on the database, that is:

**Example:** (Transaction  $T$ ) Transferring 5000\$ from account A to account B.

```
read(A);  
A := A - 5000;  
write(A);  
read(B);  
B := B + 5000;  
write(B);  
Commit;
```

Notes:

- **read(X)** transfers data item X from the database to the transaction's workspace; **write(X)** writes data item X from the workspace back to the database
- Transaction **Commit**: The transaction is completed successfully, and the database enters a new state.
- Transaction **Rollback**: The transaction is aborted, and all changes are undone, restoring the database to its previous state.

# Transaction Management

---

A transaction is a sequence of actions on the database, that is.

## Atomic

- Transaction cannot be subdivided

## Consistent

- Transaction must take the database from one valid state to another
- e.g., database is consistent with the integrity rules before and after the transaction

## Isolated

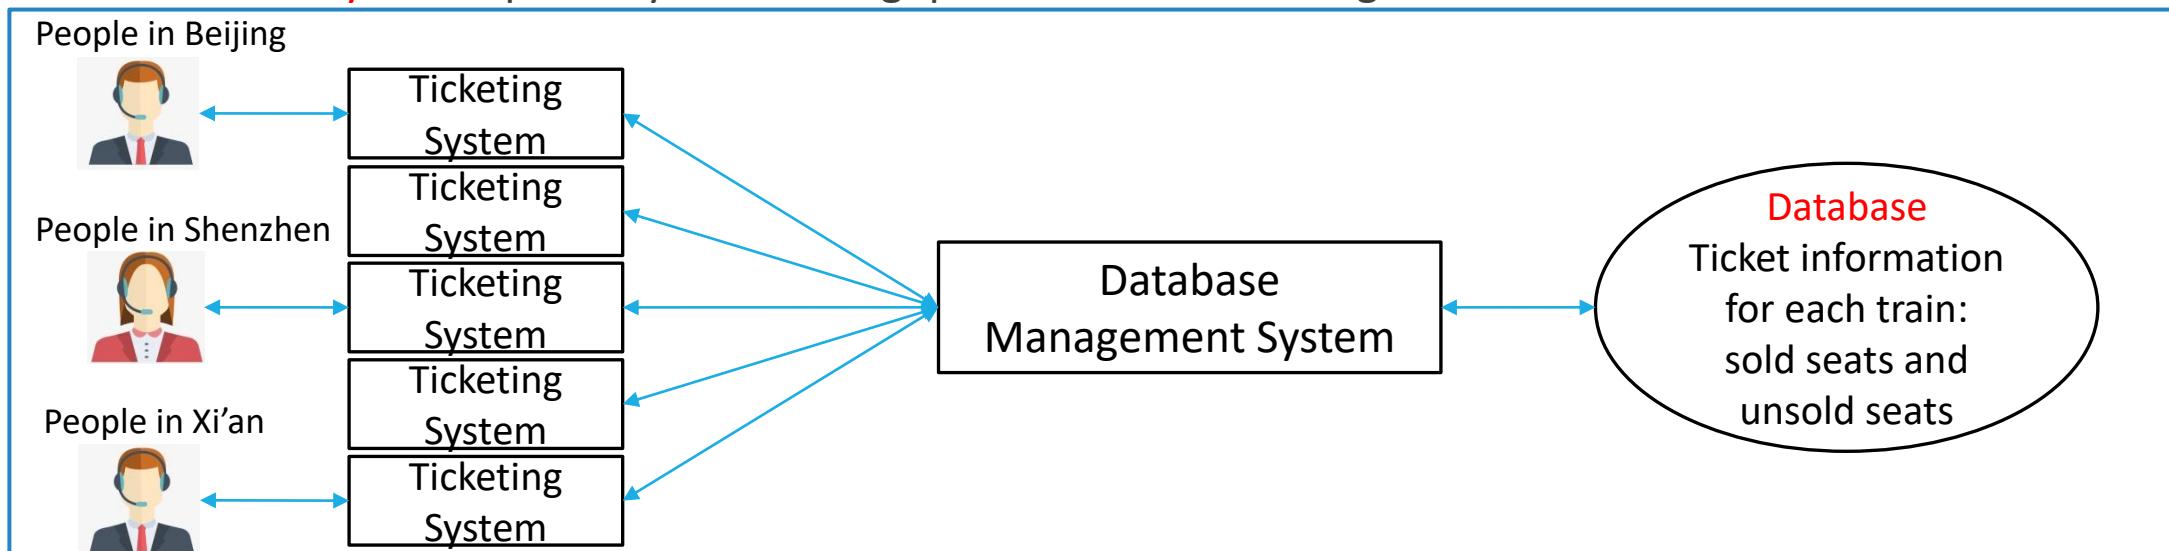
- One transaction must not interfere with another

## Durable

- Database changes are permanent. Committed transactions must remain permanent.

# Concurrency Control

- ❖ **Why is concurrency control necessary** - Multiple users executing transactions simultaneously on the database ensures no conflicts or data inconsistencies occur
  - ❖ **Train ticketing systems**: prevent overbooking and seat conflicts; ensure correct handling of ticket purchases and refunds, and accurate pricing.
  - ❖ **Online shopping platforms, course selection systems**.
  - ❖ **Concurrency** can improve system throughput and reduce waiting time.



# Transaction Scheduling and Serializability

- ❖ **Transaction scheduling:** The execution order of the basic steps (read, write, and other control operations such as locking and unlocking) of a group of transactions is called a schedule for that group of transactions.
- ❖ **Concurrent (or parallel) scheduling:** Executing the operations in a group of concurrently running transactions in a **certain order**.

| S  | T <sub>1</sub> | T <sub>2</sub> |
|----|----------------|----------------|
| 1  | Read A         |                |
| 2  | A=A-10         |                |
| 3  | Write A        |                |
| 4  | Read B         |                |
| 5  | B=B+10         |                |
| 6  | Write B        |                |
| 7  |                | Read B         |
| 8  |                | B=B-20         |
| 9  |                | Write B        |
| 10 |                | Read C         |
| 11 |                | C=C+20         |
|    |                | Write C        |

Serial  
scheduling

| S  | T <sub>1</sub> | T <sub>2</sub> |
|----|----------------|----------------|
| 1  | Read A         |                |
| 2  |                | Read B         |
| 3  | A=A-10         |                |
| 4  |                | B=B-20         |
| 5  | Write A        |                |
| 6  |                | Write B        |
| 7  | Read B         |                |
| 8  |                | Read C         |
| 9  | B=B+10         |                |
| 10 |                | C=C+20         |
| 11 | Write B        |                |
| 12 |                | Write C        |

Concurrent  
scheduling

| S  | T <sub>1</sub> | T <sub>2</sub> |
|----|----------------|----------------|
| 1  | Read A         |                |
| 2  | A=A-10         |                |
| 3  |                | Read B         |
| 4  | Write A        |                |
| 5  |                | B=B-20         |
| 6  | Read B         |                |
| 7  |                | Write B        |
| 8  | B=B+10         |                |
| 9  |                | Read C         |
| 10 | Write B        |                |
| 11 |                | C=C+20         |
| 12 |                | Write C        |

# Transaction Scheduling and Serializability

---

- ❖ A Simple **Notation** Model for Transaction Scheduling
  - ❖  $r_T(A)$ : Transaction T **reads** database object A.
  - ❖  $w_T(A)$ : Transaction T **writes** to database object A.
- ❖ Here, A and B represent **database objects**, such as tuple values or attribute values.
- ❖ This model focuses solely on the read and write operations; other computational steps that may occur in memory between these operations are not represented.

# Transaction Scheduling and Serializability

---

- ❖ **Concurrent Schedules**: A schedule is correct if and only if the resulting database state produced by the concurrent schedule is **identical** to the state produced by running those transactions serially in some order.
- ❖ **Serializability**: A schedule is considered **serializable** (or possesses the property of **serializability**) if its effect on the final state of the database is equivalent to the effect of some serial schedule, regardless of the initial state of the database.

# Transaction Scheduling and Serializability

---

- ❖ **Conflict**: Two consecutive operations in a schedule, from different transactions, that access the same data object and **at least one of them is a write operation**.
- ❖ **Conflicting** operations are **non-commutative**. Swapping their order may alter the final outcome of the transactions.  
 $A = 100$ . T1: Read(A), T2: Write(A=0), T1: print(A)  
T2: Write(A=0), T1: Read(A), T1: print(A)
- ❖ **Non-conflicting** operations are **commutative**. Swapping their order does not change the outcome of the schedule.

# Transaction Scheduling and Serializability

---

- ❖ **Conflict**: Two consecutive operations in a schedule, from different transactions, that access the same data object and **at least one of them is a write** operation.
- ❖ Two **write** operations on the **same data object** by different transactions conflict (Write-Write conflict).
- ❖ A **read** and a **write** operation on the **same data object** by different transactions conflict (Read-Write or Write-Read conflict).

# Transaction Scheduling and Serializability

- ❖ If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.
- ❖ We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule.

- ❖ Two objects: A and B

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$



$r_1(A); w_1(A); r_2(A); \color{red}{r_1(B)}; \color{red}{w_2(A)}; w_1(B); r_2(B); w_2(B)$



$r_1(A); w_1(A); \color{red}{r_1(B)}; \color{red}{r_2(A)}; w_2(A); w_1(B); r_2(B); w_2(B)$



$r_1(A); w_1(A); r_1(B); r_2(A); \color{red}{w_1(B)}; \color{red}{w_2(A)}; r_2(B); w_2(B)$



$r_1(A); w_1(A); r_1(B); \color{red}{w_1(B)}; \color{red}{r_2(A)}; w_2(A); r_2(B); w_2(B)$



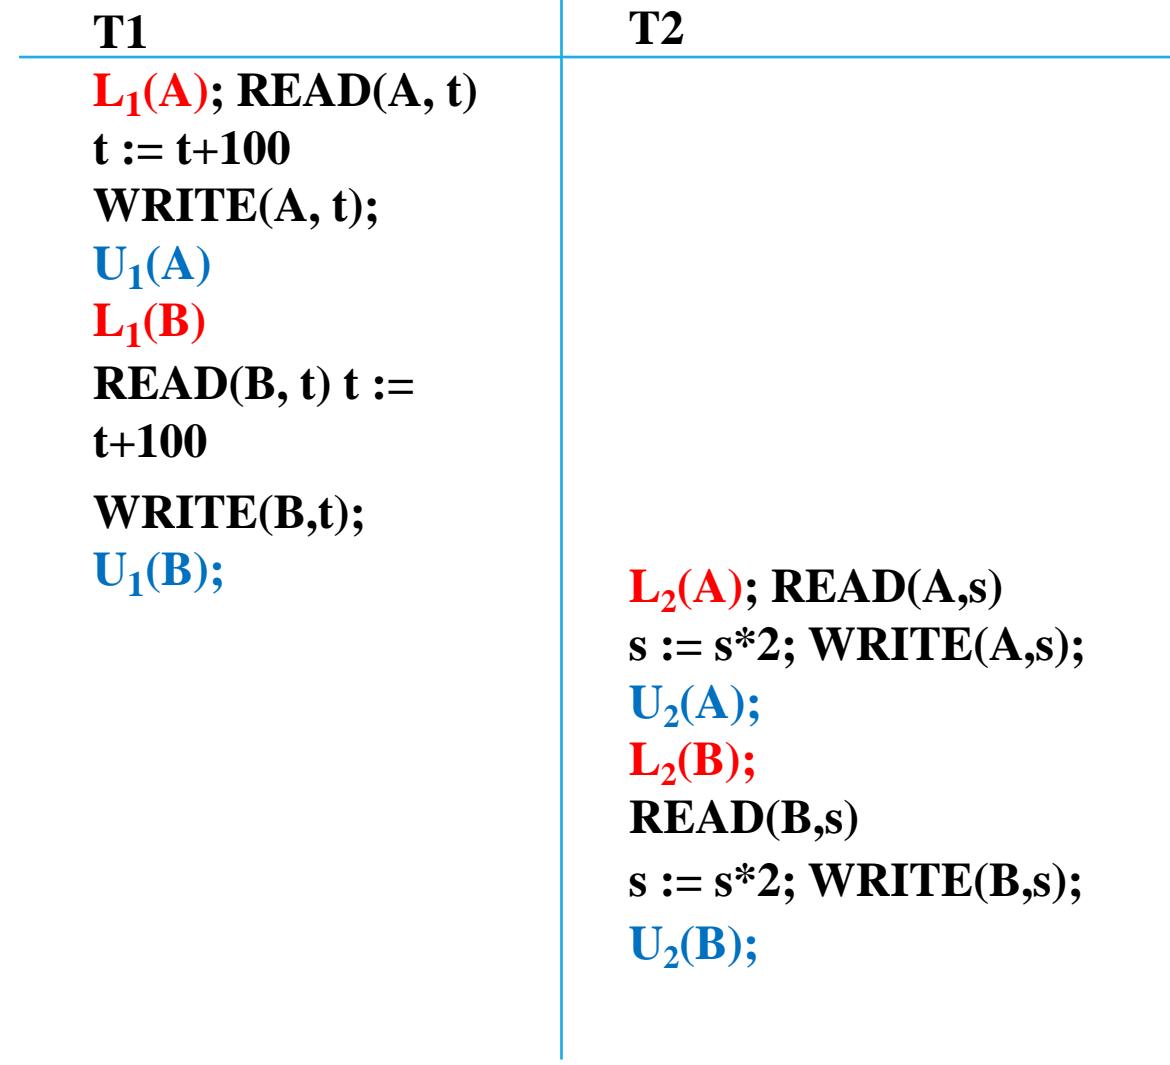
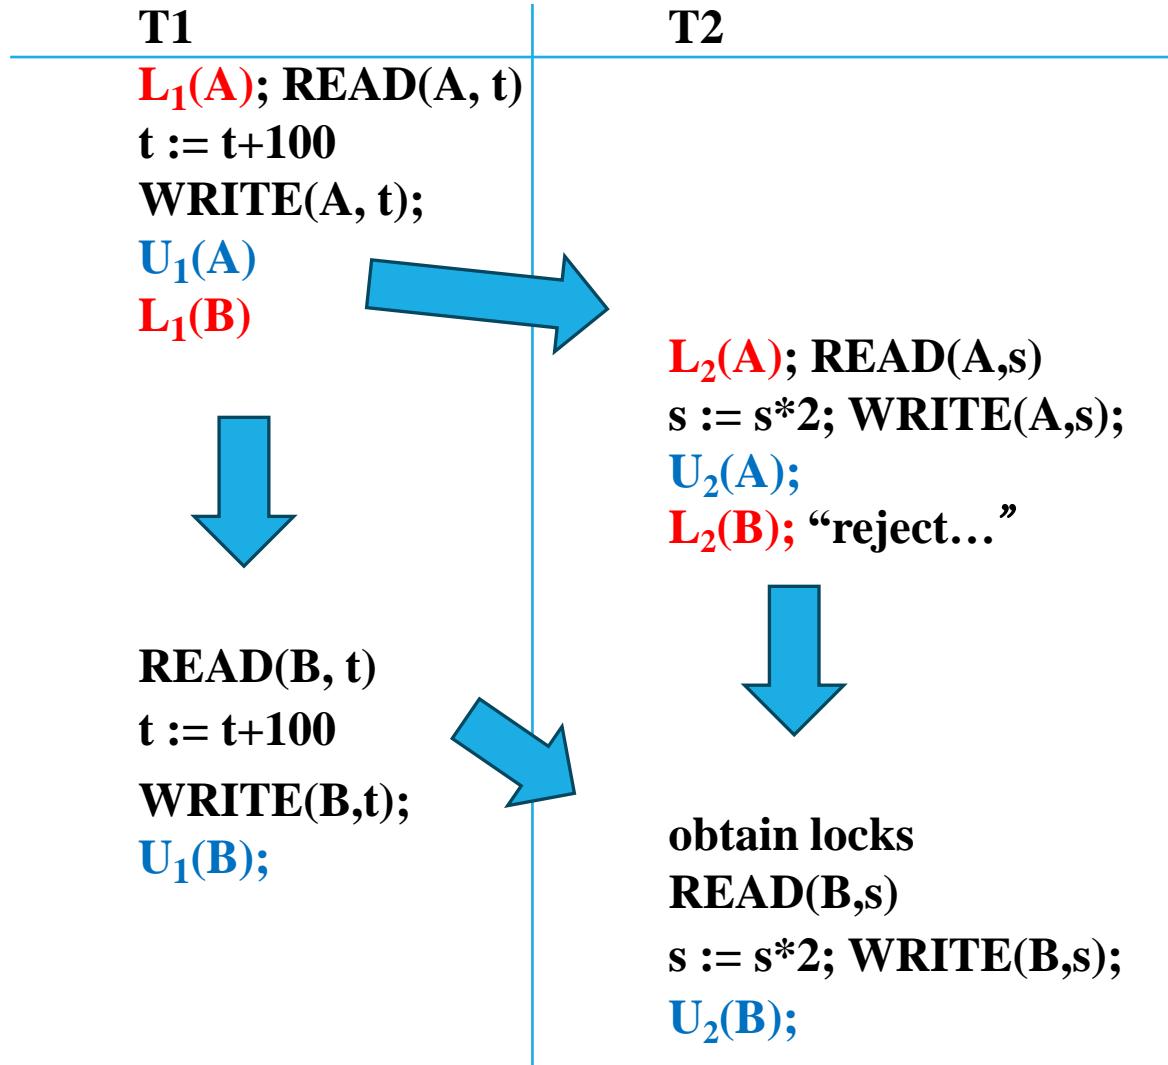
# Lock-Based Concurrency Control Protocols

---

- ❖ **Lock**: A lock is a mechanism used to control concurrency. It is a global variable that represents the "permission" to perform specific operations on a particular element.
- ❖ Before a transaction operates on an element, it must **obtain the corresponding lock**.
- ❖ If the element is already locked by another transaction and cannot be re-locked, the transaction must wait for the other transaction to release the lock.
- ❖ Once the operation is completed, the lock must be **released**. After the transaction ends, all locks held by it are released.

# Example (Lock)

The scheduler can use locks to achieve (but **does not guarantee** to achieve) conflict serializability.



# Lock-Based Concurrency Control Protocols

---

## ❖ Lock Types in Locking Protocols

- ❖ **Exclusive Lock (X):** The transaction holding the lock can read and write to the element, and no other transaction can acquire a lock on it.
  
- ❖ **Shared Lock (S):** The transaction holding the lock can only read the element, not write to it. Multiple shared locks can coexist on the same element.
  
- ❖ **Update Lock (U):** Initially a shared lock for reading, which can later be upgraded to an exclusive lock for writing.

# Lock-Based Concurrency Control Protocols

---

## ❖ Lock-compatibility matrix

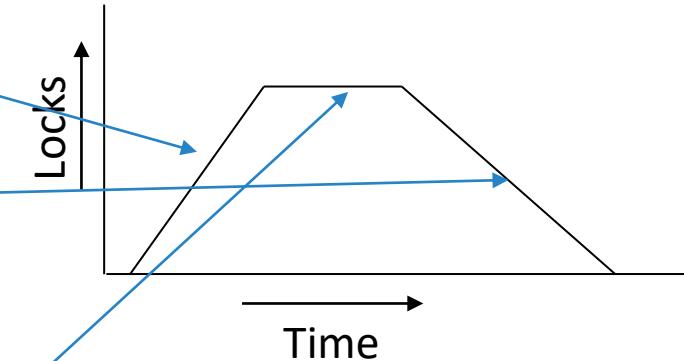
|   | S     | X     |
|---|-------|-------|
| S | true  | false |
| X | false | false |

- ❖ A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- ❖ Any number of transactions can **hold shared locks** on an item
- ❖ But if any transaction holds an exclusive on the item **no other transaction** may hold any lock on the item.

# The Two Phase Locking Protocol

---

- ❖ A protocol which ensures conflict-serializable schedules.
- ❖ Phase 1: **Growing Phase**
  - ❖ Transaction may obtain locks
  - ❖ Transaction may not release locks
- ❖ Phase 2: **Shrinking Phase**
  - ❖ Transaction may release locks
  - ❖ Transaction may not obtain locks
- ❖ The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).

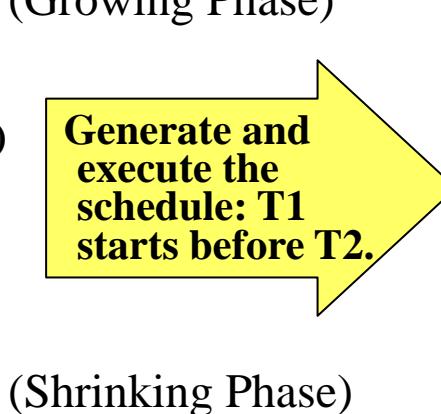


# Example (2PL)

| T1             | T2             |
|----------------|----------------|
| $R_1(B)$       | $R_2(B)$       |
| $B := B - 50;$ | $R_2(A)$       |
| $W_1(B)$       | $Display(A+B)$ |
| $R_1(A)$       |                |
| $A := A + 50;$ |                |
| $W_1(A)$       |                |

2PL

| T1             | T2             |
|----------------|----------------|
| $Lock-X(B)$    | $Lock-S(B)$    |
| $R_1(B)$       | $R_2(B)$       |
| $B := B - 50;$ | $Lock-S(A)$    |
| $W_1(B)$       | $R_2(A)$       |
| $Lock-X(A)$    | $Display(A+B)$ |
| $Unlock(B)$    | $Unlock(A)$    |
| $R_1(A)$       | $Unlock(B)$    |
| $A := A + 50;$ |                |
| $W_1(A)$       |                |
| $Unlock(A)$    |                |



| T1             | T2                |
|----------------|-------------------|
| $Lock-X(B);$   | $Lock-S(B): wait$ |
| $R_1(B);$      |                   |
| $B := B - 50;$ |                   |
| $W_1(B)$       |                   |
| $Lock-X(A)$    |                   |
| $Unlock(B)$    |                   |
| $R_1(A);$      | $Lock-S(B);$      |
| $A := A + 50;$ | $R_2(B)$          |
| $W_1(A)$       |                   |
| $Unlock(A)$    |                   |
| $Lock-S(A)$    |                   |
| $R_2(A);$      |                   |
| $Display(A+B)$ |                   |
| $Unlock(A);$   |                   |
| $Unlock(B)$    |                   |

# Example (2PL)

The two-phase locking protocol is a protocol that can potentially lead to **deadlock**!

| T1                      | T2                      |
|-------------------------|-------------------------|
| <b>Lock-X(B)</b>        | <b>Lock-S(A)</b>        |
| <b>R<sub>1</sub>(B)</b> | <b>Lock-S(B)</b>        |
| <b>B := B-50;</b>       | <b>R<sub>2</sub>(B)</b> |
| <b>W<sub>1</sub>(B)</b> | <b>R<sub>2</sub>(A)</b> |
| <b>Lock-X(A)</b>        | <b>Display(A+B)</b>     |
| <b>Unlock(B)</b>        | <b>Unlock(A)</b>        |
| <b>R<sub>1</sub>(A)</b> | <b>Unlock(B)</b>        |
| <b>A := A+50;</b>       |                         |
| <b>W<sub>1</sub>(A)</b> |                         |
| <b>Unlock(A)</b>        |                         |

Generate and execute the schedule: T1 starts before T2.

| T1                        | T2                       |
|---------------------------|--------------------------|
| <b>Lock-X(B);</b>         | <b>Lock-S(A)</b>         |
| <b>R<sub>1</sub>(B) ;</b> |                          |
| <b>B := B-50;</b>         |                          |
| <b>W<sub>1</sub>(B)</b>   |                          |
| <b>Lock-X(A): wait</b>    | <b>Lock-S(B): wait</b>   |
| <b>Unlock(B)</b>          |                          |
| <b>R<sub>1</sub>(A);</b>  |                          |
| <b>A := A+50;</b>         |                          |
| <b>W<sub>1</sub>(A)</b>   |                          |
| <b>Unlock(A)</b>          |                          |
|                           | <b>R<sub>2</sub>(B)</b>  |
|                           | <b>R<sub>2</sub>(A);</b> |
|                           | <b>Display(A+B)</b>      |
|                           | <b>Unlock(A);</b>        |
|                           | <b>Unlock(B)</b>         |

# Deadlock

---

- ❖ Neither T1 nor T2 can make progress — executing **Lock-S(B)** causes T2 to wait for T1 to release its lock on B, while executing **Lock-X(A)** causes T1 to wait for T2 to release its lock on A.
- ❖ Such a situation is called a **deadlock**.
- ❖ To handle a deadlock one of T1 or T2 must be **rolled back** and its locks released.

| T1                        | T2                       |
|---------------------------|--------------------------|
| <b>Lock-X(B);</b>         | <b>Lock-S(A)</b>         |
| <b>R<sub>1</sub>(B) ;</b> |                          |
| <b>B := B-50;</b>         |                          |
| <b>W<sub>1</sub>(B)</b>   |                          |
| <b>Lock-X(A): wait</b>    | <b>Lock-S(B): wait</b>   |
| <b>Unlock(B)</b>          |                          |
| <b>R<sub>1</sub>(A);</b>  | <b>R<sub>2</sub>(B)</b>  |
| <b>A := A+50;</b>         | <b>R<sub>2</sub>(A);</b> |
| <b>W<sub>1</sub>(A)</b>   | <b>Display(A+B)</b>      |
| <b>Unlock(A)</b>          | <b>Unlock(A);</b>        |
| <b>Unlock(B)</b>          | <b>Unlock(B)</b>         |

# Exercise

---

- Add lock and unlock instructions to transactions  $T_1$  and  $T_2$  so that they observe the two-phase locking protocol.
- Then consider: can the execution of these transactions result in a **deadlock**?

$T_1$ :

---

read(A)

---

read(B)

---

if A = 0

then B := B + 1

write(B)

---

---

$T_2$ :

---

read(B)

---

read(A)

---

if B = 0

then A := A + 1

write(A)

---

---

# Exercise

---

- ❑ Add lock and unlock instructions to transactions  $T_1$  and  $T_2$  so that they observe the two-phase locking protocol.
- ❑ Then consider: can the execution of these transactions result in a **deadlock**?

$T_1$ :

**Lock-S(A)**

read(A)

**Lock-X(B)**

read(B)

if A = 0

then B := B + 1

write(B)

**Unlock(B)**

**Unlock(A)**

$T_2$ :

**Lock-S(B)**

read(B)

**Lock-X(A)**

read(A)

if B = 0

then A := A + 1

write(A)

**Unlock(A)**

**Unlock(B)**

# Exercise

---

- ❑ Add lock and unlock instructions to transactions  $T_1$  and  $T_2$  so that they observe the two-phase locking protocol.
- ❑ Then consider: can the execution of these transactions result in a **deadlock**?

$T_1$ :

**Lock-S(A)**

read(A)

**Lock-X(B)**

read(B)

if A = 0

then B := B + 1

write(B)

**Unlock(B)**

**Unlock(A)**

$T_2$ :

**Lock-X(A)**

**Lock-S(B)**

read(B)

read(A)

if B = 0

then A := A + 1

write(A)

**Unlock(A)**

**Unlock(B)**