

Database Management

COMP 3010E FALL 2025

LECTURE 4 RELATIONAL MODEL AND LOGICAL DESIGN


Agenda

- Relational Model
- Logical Design: converting EERD to relational schema

Steps of Data Modeling

Conceptual Data Model

- Capture the nature of **relationships** among data/objects
- It is about understanding the **business rules**
- A symbolic model – easy to understand by **human**
- E-R/EER model



Logical Data Model

- Transform the conceptual model into a **logical representation**
- It is about creating the **database structure**
- Logically well-defined and can be understood by **computers**
- Focus of data analysts

CUSTOMER

Cust_ID	Cust_Name	Address
101	David Hills	...
102	Mary Lee	...

Physical Data Model

- Translate the logical model into **technical specifications** for storing/retrieving data
- It is about **storing data** to achieve data processing efficiency and data quality
- Database users and data analysts usually do not worry about this step



Index and file organization on disk

Logical Data Modeling

❑ Relational Data Model

- Represent data in the form of **relations** (vs. object-oriented classes or XML tags in other models)

❑ Relation

- A named, **two-dimensional table** of data
- Each column represents an **attribute** or a field
- Each row corresponds to a **record** or an instance

The diagram illustrates a relational database table named "CUSTOMER". The table has three columns: "CustID", "CustName", and "Address". The "CustID" column contains values 100, 101, 102, and 103. The "CustName" column contains values "Fred's Warehouse", "Bargain Bonanza", "Jasper's", and "Desks 'R Us". The "Address" column contains values "Greensboro, NC", "Moscow, ID", "Tallahassee, FL", and "Kettering, OH". A blue arrow points from the text "Table Name" to the top-left cell of the table. Another blue arrow points from the text "Attribute/Field" to the header cell "Address". A third blue arrow points from the text "Record/Instance" to the entire row for CustID 100. A fourth blue arrow points from the text "Attribute Value" to the value "Kettering, OH" in the "Address" column for CustID 103, which is enclosed in a blue oval.

CustID	CustName	Address
100	Fred's Warehouse	Greensboro, NC
101	Bargain Bonanza	Moscow, ID
102	Jasper's	Tallahassee, FL
103	Desks 'R Us	Kettering, OH

Relational Databases

❑ Relational database


- Represent data as a collection of **tables/relations**
- Establish relationships by the means of **common fields** in related tables

CUSTOMER

Cust_ID	Cust_Name	Address
101	David Hills	...
102	Mary Lee	...

ORDER


Order_ID	Cust_ID	Order_date
1100001	101	05/02/2012
1100002	102	05/03/2012



Relational Data Model

□ Historical Background

- First introduced in 1970 by E. F. Codd, an IBM research fellow
- Has a solid theoretical foundation (set theory)
- Widespread commercial success in 1980s
- By far the most commonly used in various database applications



□ Three Components

❖ **Data structure**

- Represent data in the form of tables (with rows and columns)

❖ **Data manipulation**

- Powerful operations (typically implemented using SQL)

❖ **Data integrity**

- Mechanisms to enforce business rules that maintain the integrity of data

Relation

□ Definition

- A named, two-dimensional table of data
- Consist of a set of **named columns** (attributes/fields) and **unnamed rows** (records/instances)
- Typically contains information about one specific **entity/relationship type**

□ Properties

- **Unique** table name
- Each entry is atomic (or **single valued**), i.e., no multivalued attributes
- Each column has a **unique** name
- Each row is **unique** (no duplicate rows)
- The **sequence/order** of columns (left to right) is insignificant
- The **sequence/order** of rows (top to bottom) is insignificant

CERTIFICATE

EmpID	CourseTitle	DateCompleted
100	SPSS	6/19/201X
100	Surveys	10/7/201X
140	Tax Acc	12/8/201X
110	Visual Basic	1/12/201X
110	C++	4/22/201X
150	SPSS	6/19/201X
150	Java	8/12/201X

Multivalued Attributes

- Remove multivalued attributes from tables

EMPLOYEE1					
<u>EmplID</u>	Name	DeptName	Salary	CourseTitle	DateCompleted
100	Margaret Simpson	Marketing	48,000	SPSS Surveys	6/19/201X 10/7/201X
140					
110					
EMPLOYEE2					
<u>EmplID</u>	Name	DeptName	Salary	CourseTitle	DateCompleted
190					
150					
—					
100	Margaret Simpson	Marketing	48,000	SPSS	6/19/201X
100	Margaret Simpson	Marketing	48,000	Surveys	10/7/201X
140	Alan Beeton	Accounting	52,000	Tax Acc	12/8/201X
110	Chris Lucero	Info Systems	43,000	Visual Basic	1/12/201X
110	Chris Lucero	Info Systems	43,000	C++	4/22/201X
190	Lorenzo Davis	Finance	55,000		
150	Susan Martin	Marketing	42,000	SPSS	6/19/201X
150	Susan Martin	Marketing	42,000	Java	8/12/201X

Relational Keys (1)

Primary Key (PK)

- An attribute or a combination of attributes used to **uniquely** identifies each row
- Contains no redundant information.
- Each relation **should** have a designated primary key
- May be the same as an entity's identifier, but there are exceptions
e.g., associative entities, weak entities, and relationships


Candidate Keys: all the keys that can serve as the PK.

EMPLOYEE1			
EmplID	Name	DeptName	Salary
100	Margaret Simpson	Marketing	48,000
140	Allen Beeton	Accounting	52,000
110	Chris Lucero	Info Systems	43,000
190	Lorenzo Davis	Finance	55,000
150	Susan Martin	Marketing	42,000

Relational Keys (2)

Composite key

- A primary key
- Consist of more than one attribute
- Commonly seen in relations of associative entities, weak entities, and relationships



Completes

EmplID	CourseTitle	DateCompleted
100	SPSS	6/19/201X
100	Surveys	10/7/201X
140	Tax Acc	12/8/201X
110	Visual Basic	1/12/201X
110	C++	4/22/201X
150	SPSS	6/19/201X
150	Java	8/12/201X

Relational Keys (3)

Foreign Key (FK)

- An attribute (possibly composite) in a relation
- Serve as the primary key of another relation
- Allow us to associate various relations/tables
- Provide an important mechanism to deal with relationships and enforce **referential integrity**

ORDER

<u>OrderID</u>	CustID	OrderDate
1100001	101	05/02/2012
1100002	102	05/03/2012

Foreign Key column

CUSTOMER

CustID	CustName	Address
101	David Hills	...
102	Mary Lee	...

Relational Keys (4)

SuperKey

- An attribute or a combination of attributes used to **uniquely** identifies each row
- May contains no redundant information
- The combination of all the columns is always a superkey (but not necessarily and candidate key)

EMPLOYEE1				
EmpID	Name	DeptName	Salary	
100	Margaret Simpson	Marketing	48,000	
140	Allen Beeton	Accounting	52,000	
110	Chris Lucero	Info Systems	43,000	
190	Lorenzo Davis	Finance	55,000	
150	Susan Martin	Marketing	42,000	

Schema (1)

□ Definition

- A description of the overall logical structure of the database
- A relational schema consists of a set of relations

□ Two Common Ways

❖ Short text statements

- Each relation is represented by its name and followed by the names of its attributes in parentheses
- Advantage: simple


❖ A graphical representation

- Each relation is represented by a name and a rectangle containing the attributes
- Advantage: better expressing referential integrity constraints

Schema (2)

Short Text Statements

- Format: RELATION1(Attribute1, Attribute2,...)
 - The **primary key** is underlined
 - The **foreign key** is indicated by a dashed underline



Example: a schema for four relations --

CUSTOMER(CustomerID, CustomerName, CustomerAddress, CustomerPostalCode)

ORDER(OrderID, OrderDate, CustomerID)


ORDER LINE(OrderID, ProductID, OrderQuantity)

PRODUCT(**ProductID**, **ProductDescription**, **ProductFinish**, **ProductStandardPrice**, **ProductLineID**)

Schema (3)

Graphical Representation

- Format: RelationName + rectangle
- The **primary key** is underlined
- The **foreign key** is indicated by a dashed underline
- Draw an **arrow** from each FK to the associated PK



Integrity Constraints

- Constraints or rules that limit acceptable values and actions, ensuring data integrity
- Major Types
 - Domain Constraints
 - Entity Integrity
 - Referential Integrity

Domain Constraints

- All of the values that appear in a column must be from the *same domain*
- **Domain definition:** domain name, description, data type, size, etc.

TABLE 4-1 Domain Definitions for INVOICE Attributes

Attribute	Domain Name	Description	Domain
CustomerID	Customer IDs	Set of all possible customer IDs	character: size 5
CustomerName	Customer Names	Set of all possible customer names	character: size 25
CustomerAddress	Customer Addresses	Set of all possible customer addresses	character: size 30
OrderID	Order IDs	Set of all possible order IDs	character: size 5
OrderDate	Order Dates	Set of all possible order dates	date: format mm/dd/yy
ProductID	Product IDs	Set of all possible product IDs	character: size 5
ProductDescription	Product Descriptions	Set of all possible product descriptions	character: size 25
ProductFinish	Product Finishes	Set of all possible product finishes	character: size 15
ProductStandardPrice	Unit Prices	Set of all possible unit prices	monetary: 6 digits
ProductLineID	Product Line IDs	Set of all possible product line IDs	integer: 3 digits
OrderedQuantity	Quantities	Set of all possible ordered quantities	integer: 3 digits

Domain Constraints

- ❑ All of the values that appear in a column must be from the *same domain*
- ❑ **Domain definition:** domain name, description, data type, size, etc.

- ❑ Domain of CustomerID: {four digit numbers starting from 1001 }
- ❑ Domain of Name: {any text strings}
- ❑ Domain of OrderDate: { all possible dates in yyyy-mm-dd format}
- ❑

- ❑ A relation is a subset of the cross-product (the set of all possible combinations) of all the domains

Entity Integrity

- Ensure every relation has a **primary key** with **valid** data values
- No primary key attribute may be **null**
- What is a null?
 - A value that may be assigned to an attribute with no or unknown applicable value
 - It is not a numeric “0”
 - It is not an empty string “ ”
 - A special marker indicates the absence of a data value

Entity integrity rule:

every primary key attribute must be non-null


DEPARTMENT

<u>DeptName</u>	Location	Fax
Sales	A building	3011912
Marketing	B building	NULL
Research	C building	3100988


Referential Integrity

❑ Referential integrity constraint

- Each **foreign key value** must either match a primary key value in another relation or be null
- Maintain data consistency among the rows of two associated relations
- Draw an arrow from each foreign key to the associated primary key




Note: if the relationship is mandatory, the foreign key cannot be null, e.g.
CustomerID in ORDER relation if each order must be placed by a customer.





Transform ERD into Relations

- Input: entity-relationship diagrams (ERD) or EER diagrams
- Output: relational schemas




Map Regular Entities (1)

- Entity-type name → relation name
- Simple attribute → relation attribute
- Composite attribute → simple-component attributes (improve data accessibility)
- Identifier attribute → primary key



Map Regular Entities (2)


- Multivalued attribute → **two** relations
 - The 1st contains all attributes except the multivalued attribute
 - The 2nd contains two attributes: the PK from the 1st relation + the multivalued attribute



EmployeeID	Skill
E001	C++
E002	Visual Basic
E002	R
E003	SQL

Map Regular Entities (3)

- Multivalued composite attribute → **two** relations with additional attributes



Multi-valued Attributes

- When you have multi-valued attributes in a table and you want to convert the table to relation ...

(a) Table with repeating groups

EmplID	Name	DeptName	Salary	CourseTitle	DateCompleted
100	Margaret Simpson	Marketing	48,000	SPSS	6/19/201X
		Surveys		Surveys	10/7/201X
140	Alan Beeton	Accounting	52,000	Tax Acc	12/8/201X
110	Chris Lucero	Info Systems			
190	Lorenzo Davis	Finance			
150	Susan Martin	Marketing			





(b) EMPLOYEE2 relation

EMPLOYEE2					
EmplID	Name	DeptName	Salary	CourseTitle	DateCompleted
100	Margaret Simpson	Marketing	48,000	SPSS	6/19/201X
100	Margaret Simpson	Marketing	48,000	Surveys	10/7/201X
140	Alan Beeton	Accounting	52,000	Tax Acc	12/8/201X
110	Chris Lucero	Info Systems	43,000	Visual Basic	1/12/201X
110	Chris Lucero	Info Systems	43,000	C++	4/22/201X
190	Lorenzo Davis	Finance	55,000		
150	Susan Martin	Marketing	42,000	SPSS	6/19/201X
150	Susan Martin	Marketing	42,000	Java	8/12/201X


Map Weak Entities

- Weak entity → a relation
 - Include the PK of the identifying relation as a FK attribute
 - PK: the partial identifier of the weak entity + the PK of the identifying relation




Map Binary 1:M Relationship

- Two entity types → **two** relations
 - Include the PK attribute(s) of the **one-side** entity/relation as a FK in the **many-side** entity/relation




Mapping Associative Entities

When no new keys are needed




Mapping Associative Entities (cont.)

When new keys are needed




Mapping Associative Entities (cont.)



Surrogate PK: a serial number or other system-assigned PK

Map Binary 1:1 Relationship

- Two entity types → **two** relations
 - Include the PK attribute(s) of the **mandatory-side** relation as a FK in the **optional-side** relation
 - The **relationship attributes** are also included in the **optional-side** relation



Answers

Player

<u>PlayerID</u>	Name	DOB	Position	<u>TeamID</u>	Not NULL
-----------------	------	-----	----------	---------------	----------

Team


<u>TeamID</u>	Name	Home	Year_Founded	<u>HeadCoach</u>	Not NULL
---------------	------	------	--------------	------------------	----------

Head Coach

<u>CoachID</u>	Name	Career_Wins	Career_Losses
----------------	------	-------------	---------------

Map Unary 1:M Relationship


- A single entity type → A **singe** relation
 - Add a (**recursive**) FK attribute to the relation, which references the PK values in the same relation



EmployeeID	Employee Name	Employee DOB	ManagerID
001	James	02/10/1990	002
002	Alice	10/20/1985	003
003	Mike	04/26/1980	null

Map Unary M:N Relationship


- The single entity type → one relation
 - Contain all attributes of the entity type
- The relationship → a **new** relation
 - The PK consists of two attributes, **both** are also FK attributes referencing the PK of the other relation



Example


<u>ItemNo</u>	<u>ItemDescription</u>	<u>ItemUnitCost</u>
001	Table Leg	10
002	Table Top	30
003	Dining Table	80
004	Dining Chair	30
005	Dining Set	200

<u>ItemNo</u>	<u>ComponentNo</u>	<u>Quantity</u>
003	001	4
003	002	1
005	003	1
005	004	4




Map Unary M:N Relationship

- The single entity type → one relation
 - Contain all attributes of the entity type
- The relationship → a **new** relation
 - The PK consists of two attributes, **both** are also FK attributes referencing the PK of the other relation





<u>UserID</u>
James
John
Alice
Mary



<u>User1</u>	<u>User2</u>	Since
James	Alice	2015
James	John	2016
Alice	Mary	2016

Map Ternary Relationship

- Three participating entity types → **three** relations
- Relationship → an **associative** relation
 - Include the PK attribute(s) of the three entities/relations as the PK/FK attributes




Why do we have to include
PTreatmentDate and
PTreatmentTime as two
components of the PK?

Map Supertype/Subtype Relationship

- Supertype → **one** relation
 - Contain all common attributes and the **subtype-discriminator** attribute
- Each subtype → an **individual** relation
 - Include the PK attribute of the supertype relation as the PK/FK attribute
 - Contain those attributes unique to the subtype only



What if the disjointness constraint is overlap rule?



EMPLOYEE				
EmployeeNumber	EmployeeName	EmployeeAddress	EmployeeDateHired	EmployeeType
001	James	123 A street	08/01/13	H
CONSULTANT				
CEmployeeNumber	ContractNumber	BillingRate		
HOURLY_EMPLOYEE				
HEmployeeNumber	HourlyRate			
001	\$35			
SALARIED_EMPLOYEE				
SEmployeeNumber	AnnualSalary	StockOption		

Map Supertype/Subtype Relationship

- subtype-discriminator: **composite** attribute for overlap rule



Delete Rules

- ❑ What happens if a referenced row is deleted?

CustomerID	CustomerName	Address
1	Jack	123 A street
2	Mark	245 B street
3	Jane	123 C ave

OrderID	Date	CustomerID
1	Oct-1	1
2	Nov-1	2
3	Jan-1	2

Delete Rules

- What happens if a referenced row is deleted?
- Restrict: don't allow delete of "parent" side if related rows exist in "dependent" side
- Cascade: Automatically delete "dependent" side rows that correspond with the "parent" side row to be deleted
- Set-to-Null: set the foreign key in the dependent side to null if deleting from the parent side

The choice of these options is based on your business rules.

Delete Rules: Option 1

- Restrict: don't allow delete of "parent" side if related rows exist in "dependent" side

Can delete?



CustomerID	CustomerName	Address
1	Jack	123 A street
2	Mark	245 B street
3	Jane	123 C Ave

OrderID	Date	CustomerID
1	Oct-1	1
2	Nov-1	2
3	Jan-1	2

Delete Rules: Option 2

- Cascade: Automatically delete “dependent” side rows that correspond with the “parent” side row to be deleted

CustomerID	CustomerName	Address
1	Jack	123 A street
2	Mark	245 B street
3	Jane	123 C Ave

OrderID	Date	CustomerID
1	Oct-1	1
2	Nov-1	2
3	Jan 1	2

Delete Rules: Option 3

- Set-to-Null: set the foreign key in the dependent side to null if deleting from the parent side

<u>CustomerID</u>	<u>CustomerName</u>	<u>Address</u>
1	Jack	123 A street
2	Mark	245 B street
3	Jane	123 C Ave

<u>OrderID</u>	<u>Date</u>	<u>CustomerID</u>
1	Oct-1	1
2	Nov-1	NULL
3	Jan-1	NULL

Summary

1. Entities become tables. Multi-valued attributes become new tables with composite PK.
2. 1:M relationship, add PK of 1 side into M side as FK
3. 1:1 add PK of mandatory side as FK into optional side.
4. For associative entity, follow 2. The FKs also become a composite PK if no surrogate key defined
5. For Unary relationship, follow 2-4. The only difference is you reduce the number of tables by 1
6. Ternary: follow 4. Also add necessary attributes into the PK.
7. Subtype: each becomes a new table. Use supertype's PK as its own PK (also FK). Don't copy Supertype's other attributes.